

OpenCL ICD Loader Status and Perspectives

Brice Videau
Argonne National Laboratory

Khronos F2F Meeting, Montréal, 2023-10-27



OpenCL ICD Loader Status

OpenCL ICD Loader Features

- Multiplexing of vendor drivers
 - Through vendor provided dispatch tables
- Provides a layering system
 - Global layers, enabled for each call
 - Introspection utility (`cllayerinfo`)
- Configuration through environment variables (or registry on windows)

OpenCL ICD Loader Issues

- Not a well behaved library
 - Leaks memory
 - Cannot be unloaded gracefully
 - Doesn't unload vendor drivers
 - Doesn't de-initialize and unload layers
- Dispatching issues
 - The loader doesn't know the size of the vendor provided dispatch table, leading to segfaults when calling outside the dispatch table range

OpenCL ICD Loader Missing Features

- No layer can be added while inside the application
 - Instance layers (Vulkan)

Potential Solutions



Library Behavior

- Expected behavior
 - libOpenCL.xx should not leak memory
 - libOpenCL.xx should be able to be opened and closed
- Solution: library destructor (and constructor?)
 - Leverages host system synchronization to avoid race conditions
- Caveats:
 - Vendor libraries need to support being loaded/unloaded (optional in vulkan, env variable) could be an extension
 - Layers need an explicit de-initialization function (new version of API)

New Dispatch Strategy - 1

- Expected behavior
 - libOpenCL.xx should not rely on a vendor provided dispatch table
- Solution: loader managed dispatch (Vulkan)
 - API to query all entry points: clGetFunctionAddressForPlatformKHR (queried through clGetExtensionFunctionAddress)
 - Change OpenCL objects layouts (more on the next slide)
 - Maintain backward compatibility?


New Dispatch Strategy - 2

- OpenCL objects layouts


- `cl_khr_icd2`
- **Currently:** `struct _cl_platform_id { cl_icd_dispatch *dispatch; };`
- **Proposed:** `struct _cl_platform_id { cl_icd_dispatch *dispatch; void *disp_data; };`
- Use `clGetPlatformIDs` entry (first entry of the dispatch table, unused by the loader) to look a for a tag.
- A well chosen tag should allow distinguishing objects from old OpenCL implementations
 - Low bits set, not a valid pointer to a function on most architectures
 - 64 bit: `0x4F50454E434C3331` (ASCII "OPENCL31")
 - 32 bit: `0x434C3331` (ASCII "CL31")
- `disp_data` can be set by the loader (but could also be set by compliant implementation, more on that later)
- `disp_data` would contain necessary dispatch information from the loader

New Dispatch Strategy - 3

- As of today ICD Loader dispatch:



- Backward compatible Strategy:



New Dispatch Strategy - 4

- Use `disp_data` to store dispatch information about object
 - Loader managed dispatch tables (akin to Vulkan)
 - Proof of concept here:
<https://github.com/Kerilk/OpenCL-ICD-Loader/commits/managed-dispatch>
 - Populated through new `clGetFunctionAddressForPlatformKHR`
 - When objects are created the `disp_data` is set to that of their parent
 - Can be done by the loader, but not in extension functions like:

```
extern CL_API_ENTRY cl_int CL_API_CALL clEnqueueSVMUnmapARM(cl_command_queue command_queue,  
void* svm_ptr, cl_uint num_events_in_wait_list, const cl_event* event_wait_list, cl_event* event);
```
 - Maybe better to let the vendor driver copy the `disp_data` pointer from the `command_queue` to the `event`

Instances and Layers



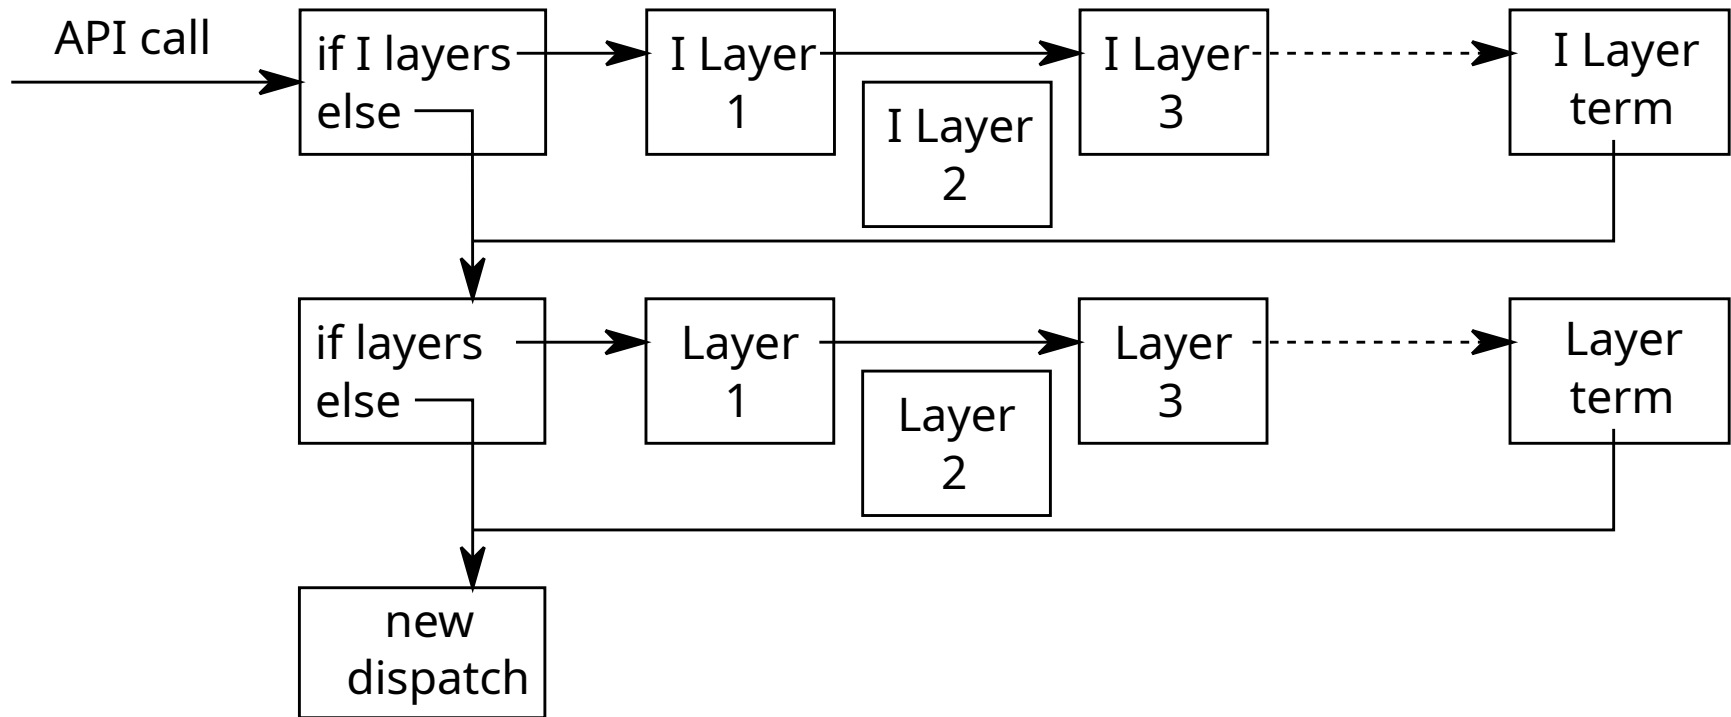
Can We Extend Further?

- Proposed solutions would fix current shortcomings
 - Library behavior
 - Dispatching issues
- Can we extend things further?
 - Application controlled
 - Instances
 - layers
 - Similar to what Vulkan proposes
 - Sketched by Ben Ashbaugh (Intel) in 2018

Yes We Can (or at least we should be able to...)

- Vulkan proved loader managed dispatch is enough for application controlled layers
- Application controlled layer without instances:
 - Would be limited to per platform, since dispatch data is inherited from platform to devices
- Idea, similar to Vulkan: create an instance to encapsulate an application controlled layer setup (and other state potentially, loader implemented)
 - `clCreateInstance(...layers configuration...)` / `clDestroyInstance` / `clGetPlatformIDsForInstance`
- Add concept of instance platforms and devices, instance specific handles for platform and devices (has to be implemented by vendors)
 - `clGetInstancePlatformsIDsKHR(cl_instance instance, cl_uint num_entries, cl_platform_id *platforms, cl_uint *num_platforms)`
 - `clReleaseInstancePlatformIDsKHR()`
 - `clGetDeviceIDs` called with an instance platform would return instance devices.
 - Instance handles can be used as regular handles in the rest of the API (like sub-devices)

Instance Layers Dispatching



Conclusion

cl_khr_icd2

- As outlined cl_khr_icd2 would allow
 - Fixing library behavior
 - Solving dispatching issues
 - Enabling new features such as instance layers
 - Maintain backward compatibility
 - Instances would only be supported for ck_khr_icd2 implementations
 - Legacy drivers would not be unloaded
 - New implementations would continue working on old loaders

Impact on Vendors

- Extension specifying if vendor driver can be unloaded? (or part of cl_khr_icd2)
- cl_khr_icd2 would be
 - 3 new entry points
 - clGetFunctionAddressForPlatformKHR
 - clGetInstancePlatformsIDsKHR
 - clReleaseInstancePlatformIDsKHR
 - 2 new object types: instance platform and devices (but not new C types)
 - A pointer copy per object creation
 - Most probably a few info queries
- Could reduce this overhead if we abandon instances

Acknowledgment

This research used resources of the Argonne Leadership Computing Facility, a U.S. Department of Energy (DOE) Office of Science user facility at Argonne National Laboratory and is based on research supported by the U.S. DOE Office of Science-Advanced Scientific Computing Research Program, under Contract No. DE-AC02-06CH11357.